

**Assignment-2**  
**Group: 16**  
**Rashi Dubey**  
**Bhaumik Icchaporia**

1. Use the read.table command to load this data into R. Make sure you set the 'header' option.

Ans:

```
>az_char<-read.table("C:/Users/dubey/OneDrive/Quarter3/Pattern/assign2/az-5000.txt",header = T)
>is.data.frame(az_char)
```

```
> is.data.frame(az_char) ##
[1] TRUE
```

Output:

	char	x1	y1	x2	y2	x3	y3	x4	y4	x5	y5	x6	y6	x7	y7	x8
1	n	0.1875	0.140625	0.09375	0.515625	0	0.8828125	0.1796875	0.5078125	0.4609375	0.140625	0.640625	0.109375	0.515625	0.5390625	0.3828125
2	h	0	0	0.046875	0.3125	0.0625	0.671875	0.15625	0.765625	0.328125	0.4609375	0.4609375	0.3359375	0.625	0.4609375	0.6015625
3	y	0.0078125	0.0859375	0.140625	0.1875	0.2578125	0.375	0.4140625	0.421875	0.5078125	0.1484375	0.484375	0.1953125	0.28125	0.5234375	0.125
4	h	0	0	0.0390625	0.25	0.0546875	0.5859375	0.046875	0.8359375	0.1015625	0.671875	0.2421875	0.625	0.3515625	0.828125	0.40625
5	y	0.03125	0	0.03125	0.2421875	0.1953125	0.25	0.3125	0.140625	0.40625	0.0390625	0.3046875	0.265625	0.203125	0.46875	0.0859375
6	q	0.453125	0.015625	0.1015625	0.0703125	0	0.3203125	0.2734375	0.3203125	0.3984375	0.1484375	0.40625	0.3671875	0.3984375	0.703125	0.328125
7	y	0.0234375	0.1640625	0.15625	0.3125	0.265625	0.40625	0.4375	0.53125	0.421875	0.1640625	0.3359375	0.3515625	0.203125	0.609375	0.1015625
8	t	0.0546875	0	0.203125	0.1328125	0.1953125	0.296875	0.15625	0.5703125	0.1484375	0.7421875	0.203125	0.9453125	0.3359375	0.9921875	0.0625
9	u	0	0.0234375	0.0078125	0.359375	0.0078125	0.65625	0.1171875	0.96875	0.40625	0.7265625	0.5625	0.421875	0.546875	0.3515625	0.53125
10	q	0.328125	0.1015625	0.0703125	0.0546875	0	0.3359375	0.21875	0.3359375	0.296875	0.1953125	0.296875	0.546875	0.296875	0.9375	0.484375
11	w	0.125	0.3359375	0	0.703125	0.1640625	0.9765625	0.4296875	0.671875	0.5703125	0.40625	0.4375	0.6953125	0.6328125	0.9765625	0.90625
12	c	0.6953125	0.2734375	0.9296875	0.109375	0.7890625	0.1875	0.546875	0.3671875	0.203125	0.5859375	0.0078125	0.765625	0.0546875	0.9921875	0.4609375
13	t	0.34375	0.078125	0.3359375	0.234375	0.296875	0.390625	0.2578125	0.6328125	0.2265625	0.796875	0.296875	0.9921875	0.4921875	0.9609375	0.1484375
14	o	0.4296875	0.046875	0.0859375	0.09375	0	0.4453125	0.109375	0.8359375	0.2890625	0.9921875	0.625	0.953125	0.859375	0.671875	0.875
15	m	0	0.609375	0.078125	0.8125	0.1953125	0.96875	0.3203125	0.7734375	0.4765625	0.8203125	0.6171875	0.84375	0.7109375	0.703125	0.9140625
16	b	0.046875	0.03125	0.09375	0.4453125	0.09375	0.8515625	0	0.828125	0.078125	0.4921875	0.328125	0.421875	0.546875	0.6953125	0.375
17	b	0.0625	0	0.046875	0.3515625	0.0078125	0.671875	0.015625	0.7109375	0.1484375	0.4609375	0.375	0.46875	0.3515625	0.734375	0.203125
18	j	0.3046875	0.421875	0.296875	0.5234375	0.3125	0.671875	0.3359375	0.7734375	0.3359375	0.875	0.2578125	0.953125	0.1484375	0.9765625	0.046875
19	m	0.0078125	0.1484375	0.046875	0.484375	0	0.90625	0.2578125	0.234375	0.5078125	0.3984375	0.5703125	0.7109375	0.609375	0.3203125	0.90625
20	l	0.03125	0.03125	0.0234375	0.140625	0.0234375	0.2421875	0.015625	0.3671875	0.015625	0.5546875	0.0078125	0.671875	0.0078125	0.7734375	0
21	v	0	0.3203125	0.0390625	0.1328125	0.140625	0.265625	0.140625	0.6015625	0.03125	0.890625	0.140625	0.921875	0.3515625	0.7265625	0.7265625
22	r	0	0.0859375	0.0390625	0.1953125	0.09375	0.3828125	0.1171875	0.546875	0.140625	0.75	0.1328125	0.9453125	0.140625	0.1953125	0.2109375
23	f	0.4453125	0.0703125	0.2734375	0.0703125	0.1875	0.328125	0.1796875	0.5703125	0.171875	0.8984375	0.1171875	0.9140625	0.0390625	0.7578125	0.1640625
24	n	0	0	0.03125	0.3203125	0.0859375	0.6875	0.2421875	0.5703125	0.34375	0.109375	0.5390625	0.0234375	0.7265625	0.3671875	0.859375
25	j	0.1875	0.2578125	0.1796875	0.40625	0.2265625	0.5625	0.234375	0.625	0.2421875	0.796875	0.1953125	0.96875	0.09375	0.9921875	0.0078125
26	a	0.4140625	0.7109375	0.359375	0.2421875	0.0234375	0.53125	0.046875	0.875	0.3828125	0.859375	0.5078125	0.6328125	0.3828125	0.609375	0.625
27	b	0.390625	0.109375	0.3515625	0.2734375	0.359375	0.671875	0.3984375	0.875	0.5	0.625	0.75	0.625	0.8828125	0.828125	0.4765625
28	u	0.0390625	0.0859375	0	0.3046875	0.015625	0.6640625	0.296875	0.8125	0.640625	0.6953125	0.7421875	0.4453125	0.71875	0.234375	0.765625
29	m	0	0.546875	0.203125	0.359375	0.234375	0.7734375	0.28125	0.7265625	0.46875	0.46875	0.546875	0.859375	0.71875	0.6640625	0.9453125
30	s	0.7109375	0.09375	0.5078125	0	0.2578125	0.1015625	0.3125	0.4140625	0.5	0.5390625	0.7734375	0.765625	0.671875	0.9921875	0.3359375
31	k	0.0859375	0.0859375	0.0625	0.328125	0.0390625	0.6640625	0.015625	0.96875	0.15625	0.28125	0.2578125	0.2265625	0.078125	0.40625	0.1796875
32	h	0.78125	0.4609375	0.2109375	0.0625	0.09375	0.359375	0.0078125	0.65625	0.0703125	0.8984375	0.4140625	0.6171875	0.5078125	0.7109375	0.4453125

- 1.b) Use the sample command to randomly select 80% of the data for training.

Ans:

```
>indices <- sample(1:nrow(az_char), size = (0.8*nrow(az_char)))##sampling 80% of the data
>training <- az_char[indices,] ##putting the sampled data in training vector
>test <- az_char[-indices,] ##putting the rest dataset in the test_data
```

Output:

	char	x1	y1	x2	y2	x3	y3	x4	y4	x5	y5	x6	y6	x7	y7	x8	y8	x9	y9
4671	u	0.0000000	0.3046875	0.0781250	0.6328125	0.1875000	0.8984375	0.5703125	0.9609375	0.9218750	0.6406250	0.9375000	0.2734375	0.8593750	0.1953125	0.9296875	0.4765625	0.9765625	0.9375000
4143	y	0.0078125	0.0000000	0.1562500	0.1640625	0.3046875	0.3906250	0.7812500	0.2734375	0.6250000	0.4687500	0.5000000	0.6093750	0.3671875	0.7500000	0.2734375	0.8515625	0.1093750	0.9921875
4501	j	0.3671875	0.0156250	0.3906250	0.4140625	0.4140625	0.5703125	0.4062500	0.7500000	0.3671875	0.9062500	0.2734375	0.9843750	0.1093750	0.9609375	0.0703125	0.9375000	0.1718750	0.9687500
1792	t	0.2343750	0.0000000	0.1796875	0.2031250	0.1562500	0.4218750	0.1328125	0.7343750	0.1875000	0.9921875	0.3671875	0.8984375	0.4531250	0.7578125	0.1015625	0.4453125	0.3750000	0.3593750
3037	q	0.3593750	0.0156250	0.1562500	0.0390625	0.0000000	0.1796875	0.0703125	0.3671875	0.3046875	0.1171875	0.3125000	0.3203125	0.2890625	0.5859375	0.2968750	0.8281250	0.3281250	0.9921875
1686	o	0.8046875	0.0625000	0.3906250	0.0625000	0.0468750	0.2890625	0.0000000	0.6171875	0.1406250	0.9375000	0.4531250	0.9921875	0.8046875	0.8906250	0.9843750	0.6494375	0.8984375	0.2500000
2155	b	0.0234375	0.0078125	0.0390625	0.3359375	0.0390625	0.6406250	0.0390625	0.9531250	0.0781250	0.7187500	0.2109375	0.5703125	0.4375000	0.7265625	0.2500000	0.9218750	0.0078125	0.8437500
3821	x	0.7265625	0.1250000	0.6796875	0.1796875	0.4609375	0.4609375	0.2500000	0.6718750	0.1250000	0.9218750	0.1953125	0.1953125	0.4531250	0.4296875	0.7187500	0.6171875	0.9531250	0.8046875
3309	y	0.0078125	0.0937500	0.1640625	0.2343750	0.2890625	0.3906250	0.4609375	0.0390625	0.3828125	0.2343750	0.2968750	0.4218750	0.2187500	0.6093750	0.1484375	0.7968750	0.0625000	0.9843750
1890	k	0.0000000	0.2343750	0.1875000	0.2812500	0.2109375	0.7187500	0.3906250	0.7500000	0.7031250	0.4453125	0.7890625	0.3125000	0.5468750	0.5000000	0.6718750	0.9140625	0.9843750	0.9921875
132	r	0.0000000	0.1406250	0.0468750	0.2968750	0.0781250	0.5468750	0.1015625	0.7109375	0.1015625	0.9375000	0.0390625	0.1250000	0.2187500	0.0000000	0.4296875	0.0078125	0.5546875	0.1250000
4506	b	0.1093750	0.0000000	0.1250000	0.2578125	0.1406250	0.5937500	0.1718750	0.7968750	0.2500000	0.5468750	0.4687500	0.5937500	0.4765625	0.8359375	0.2109375	0.9609375	0.0156250	0.9921875
942	k	0.0078125	0.0156250	0.0625000	0.5000000	0.1406250	0.8906250	0.1562500	0.7343750	0.3281250	0.4218750	0.4843750	0.1875000	0.2890625	0.4843750	0.4843750	0.7265625	0.6250000	0.9531250
4415	y	0.0312500	0.0781250	0.1796875	0.1953125	0.2812500	0.3203125	0.4453125	0.0468750	0.3906250	0.2578125	0.2968750	0.4765625	0.1875000	0.6718750	0.0937500	0.8281250	0.0000000	0.9921875
1775	t	0.6875000	0.0156250	0.5000000	0.0000000	0.4531250	0.3125000	0.3828125	0.5312500	0.2812500	0.9140625	0.5390625	0.9531250	0.0000000	0.3750000	0.2734375	0.3593750	0.5078125	0.3203125
2838	h	0.0078125	0.0000000	0.0000000	0.2187500	0.0000000	0.4296875	0.0468750	0.6250000	0.1484375	0.7265625	0.2421875	0.5312500	0.3750000	0.6328125	0.4218750	0.8046875	0.4453125	0.9687500
4995	m	0.0468750	0.2656250	0.0859375	0.6406250	0.1718750	0.7031250	0.2812500	0.1718750	0.5078125	0.5468750	0.5625000	0.8203125	0.7031250	0.2656250	0.9375000	0.4531250	0.9843750	0.9921875
2757	d	0.3515625	0.6875000	0.1640625	0.5156250	0.0000000	0.8281250	0.3359375	0.9218750	0.5078125	0.5625000	0.6015625	0.1718750	0.5468750	0.2734375	0.4765625	0.7500000	0.8984375	0.8281250
2001	u	0.0156250	0.0000000	0.0000000	0.4765625	0.2031250	0.8984375	0.5781250	0.8046875	0.5703125	0.3203125	0.5390625	0.1640625	0.5859375	0.6406250	0.7109375	0.9375000	0.8046875	0.6875000
4572	o	0.6640625	0.3593750	0.3750000	0.2343750	0.0703125	0.4062500	0.0312500	0.7187500	0.3750000	0.9531250	0.7968750	0.9296875	0.9843750	0.6875000	0.7890625	0.4062500	0.4140625	0.3046875
1929	i	0.0390625	0.0000000	0.0000000	0.1406250	0.0000000	0.3125000	0.0312500	0.4531250	0.0781250	0.5937500	0.1328125	0.7265625	0.2031250	0.8281250	0.2421875	0.8906250	0.3281250	0.9921875
4419	z	0.2343750	0.2187500	0.4140625	0.2109375	0.8984375	0.1796875	0.8125000	0.4296875	0.3984375	0.6328125	0.1171875	0.1875000	0.1562500	0.9921875	0.6171875	0.9062500	0.9375000	0.8671875
4507	e	0.2343750	0.1328125	0.2812500	0.2812500	0.0390625	0.2578125	0.3593750	0.9531250	0.6406250	0.9140625	0.5390625	0.9531250	0.3750000	0.2734375	0.3593750	0.5078125	0.3203125	0.9921875

training 4000 obs. of 19 variables

Shows the total of 4000 observations of 19 variables.

1.c) Use the table command to show the number of cases per class in the training data.

Ans:

```
>case_by_class <- table(training$char)
>print(case_by_class)
```

Output

```
 a  b  c  d  e  f  g  h  i  j  k  l  m  n  o  p  q  r  s  t  u  v  w  x  y  z
140 136 159 155 160 135 147 146 154 156 133 178 153 131 176 165 167 169 164 153 178 165 152 152 156 120
```

2. Linear Discriminant Analysis

2.a) Use the c() command to create a vector of prior probabilities equal to 1/26 for each class.

Ans:

```
>prior_vector <- c(rep(1/26, each=26))
>print(prior_vector)
```

Output: Each of the 26 values having a prior probability of **0.03846154**.

```
> print(prior_vector)
[1] 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154
[13] 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154 0.03846154
[25] 0.03846154 0.03846154
> |
```

2.b) Use the lda command to run linear discriminant analysis on the training data with the equal priors above. You may need to load the "MASS" package. In R, the syntax "char ~." indicates the formula for our functional model – i.e., that we are trying to predict char (column one in the data) as a function of all the other variables.



Ans.

```
>require(MASS)
>lda_model <- lda(formula = char~., data = training, prior = prior_vector) ##using lda
to on training dataset with each value of equal prior probability to predict character
>print(lda_model)
```

Output:

Name	Type	Value
lda_model	list [10] (S3: lda)	List of length 10
prior	double [26]	0.0385 0.0385 0.0385 0.0385 0.0385 0.0385 ...
counts	integer [26]	140 144 165 145 155 151 ...
means	double [26 x 18]	0.4106 0.0936 0.6551 0.3735 0.1537 0.2893 0.3120 0.1183 0.1744 0.4880 0.4342 0.3 ...
scaling	double [18 x 18]	3.1638 1.0381 -0.5445 0.8042 -0.6945 -1.6957 2.3666 0.1841 -0.7108 -1.0464 ...
lev	character [26]	'a' 'b' 'c' 'd' 'e' 'f' ...
svd	double [18]	24.4 23.3 16.9 14.6 14.4 13.3 ...
N	integer [1]	4000
call	language	lda(formula = char ~ ., data = training, prior = prior_vector)
terms	formula	char ~ x1 + y1 + x2 + y2 + x3 + y3 + x4 + y4 + x5 + y5 + x6 + y6 + x7 + y7 ...
xlevels	list [0]	List of length 0

2.c) c. Combine the functions table and predict to print a “confusion” matrix on the test data. This is a 26x26 matrix with diagonal elements equal to correct classifications and off-diagonal elements equal to mistakes. Which character had the best/worst performance?

Ans:

# prediction

```
>test <- az_char[-indices] ## 20% data put in test vector
>predict_lda_test <- predict(object = lda_model, newdata = test)
>conf_matrix_test <- table(test$char, predict_lda_test$class)
>print(conf_matrix_test)
```

Output:

```

> require(MASS)
Loading required package: MASS
> lda_model <- lda(formula = char~.,
+                   data = training,
+                   prior = prior_vector)
> # prediction
> predict_lda_test <- predict(object = lda_model, newdata = test)
> conf_matrix_test <- table(test$char, predict_lda_test$class)
> print(conf_matrix_test)

```

	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z	
a	21	0	0	4	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	4	1	7	
b	0	40	0	0	2	1	0	1	0	0	0	2	0	0	0	0	0	0	0	1	0	0	0	0	0	1	0
c	1	0	33	1	3	0	0	1	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0
d	0	0	0	25	0	0	0	1	0	0	0	0	0	0	0	0	0	0	1	1	0	0	0	4	0	0	0
e	0	0	0	0	32	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	1	
f	0	0	1	0	0	37	0	0	0	0	0	2	0	0	0	5	0	2	0	4	0	0	0	1	0	0	0
g	0	0	0	0	1	2	22	0	0	0	0	0	0	0	0	0	4	0	3	1	0	0	0	2	10	0	0
h	3	0	0	0	3	0	0	23	0	0	2	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0
i	0	0	0	0	0	1	0	0	28	6	0	1	0	0	0	0	0	0	0	2	0	0	0	0	0	0	0
j	0	2	0	0	0	1	0	0	5	25	0	0	0	0	0	0	0	0	0	2	0	0	0	1	2	0	0
k	1	0	1	0	0	0	0	4	0	0	26	2	1	1	0	0	0	0	0	0	4	0	0	0	0	0	0
l	0	0	0	1	1	2	0	1	0	0	0	19	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
m	1	0	0	1	0	0	0	0	0	0	0	0	31	2	0	0	0	0	0	0	4	0	0	0	0	0	0
n	0	0	0	0	0	0	0	1	0	0	1	0	0	26	0	0	1	0	0	0	4	0	0	0	0	0	0
o	0	0	1	0	0	0	0	0	0	0	0	0	0	0	25	0	0	1	1	1	0	0	0	6	0	0	0
p	0	0	0	0	0	2	0	0	0	0	0	0	0	0	0	38	0	0	0	0	0	0	2	0	1	0	0
q	0	0	0	0	0	1	2	0	0	0	0	1	0	0	0	0	23	0	1	0	0	0	0	1	1	0	0
r	0	0	0	0	0	0	0	0	0	0	0	0	0	2	0	0	3	36	0	1	1	3	0	0	1	1	0
s	0	0	0	0	0	1	1	0	0	0	0	0	0	1	0	0	0	0	39	0	0	0	0	0	1	0	0
t	0	0	0	0	2	2	0	0	1	1	0	1	0	0	0	0	0	0	0	27	1	0	0	2	0	1	0
u	2	0	0	0	0	0	0	1	0	0	0	0	1	1	1	0	0	0	0	1	33	1	1	0	0	0	0
v	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	3	29	2	0	0	0	0
w	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	2	44	0	0	0	0
x	1	1	0	1	0	0	0	0	0	0	0	0	0	0	2	0	0	2	0	0	2	2	0	17	3	0	0
y	0	0	0	0	0	2	0	0	0	0	0	1	0	0	0	0	1	1	1	1	0	0	0	0	28	0	0
z	0	0	0	0	0	0	0	0	0	1	0	1	0	0	0	0	1	0	0	3	0	0	0	0	1	33	0

```

> |

```

```
> print(conf_matrix_test)
```

	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z
a	21	0	0	4	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	3	1	8
b	0	40	0	0	2	1	0	1	0	0	0	2	0	0	0	0	0	0	0	1	0	0	0	0	1	0
c	0	0	34	1	3	0	0	1	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0
d	0	0	0	25	0	0	0	1	0	0	0	0	0	0	0	0	0	0	1	1	0	0	0	4	0	0
e	0	0	0	0	33	1	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0
f	0	0	0	1	0	37	0	0	0	0	0	2	0	0	0	6	0	1	0	4	0	0	0	1	0	0
g	0	0	0	0	1	2	24	0	0	0	0	0	0	0	0	0	4	0	4	1	0	0	0	2	7	0
h	3	0	0	0	3	0	0	23	0	0	2	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
i	0	0	0	0	0	1	0	0	31	5	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
j	0	2	0	0	0	1	0	0	5	25	0	0	0	0	0	0	0	0	0	2	0	0	0	1	2	0
k	1	0	1	0	0	0	0	5	0	0	26	1	1	1	0	0	0	0	0	0	4	0	0	0	0	0
l	0	0	0	1	1	2	0	1	0	0	0	19	0	0	0	0	0	0	0	0	0	0	0	0	0	0
m	1	0	0	0	0	0	0	0	0	0	0	0	32	2	0	0	0	0	0	0	4	0	0	0	0	0
n	0	0	0	0	0	0	0	1	0	0	1	0	0	25	0	0	1	0	0	0	5	0	0	0	0	0
o	0	0	1	0	0	0	0	0	0	0	0	0	0	0	25	0	0	1	1	1	0	0	0	6	0	0
p	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	39	0	0	0	0	0	1	0	1	1	0
q	0	0	0	0	0	0	2	0	0	0	0	1	0	0	0	0	24	0	1	0	0	0	0	0	2	0
r	0	0	0	0	0	0	0	0	0	0	1	0	0	1	0	0	3	36	0	1	1	3	0	0	1	1
s	0	0	0	0	0	1	1	0	0	0	0	1	0	0	0	0	0	0	39	0	0	0	0	0	1	0
t	0	0	0	0	2	2	0	0	1	1	0	1	0	0	0	0	1	0	26	1	0	0	0	2	0	1
u	2	0	0	0	0	0	0	1	0	0	0	0	1	1	1	0	0	0	0	1	33	1	1	0	0	0
v	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	3	29	2	0	0	0	0
w	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	2	44	0	0	0	0
x	1	0	1	0	0	0	0	0	0	0	0	0	0	1	2	0	0	2	0	0	2	2	0	17	3	0
y	0	0	0	0	0	1	0	0	0	0	1	0	0	0	0	1	1	0	2	0	0	0	0	29	0	0
z	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	1	0	0	3	0	0	0	0	1	34

**Character b** has the best performance -  $40/42 = 0.952388$ .

Accuracy rate =  $0.95238 \times 100 = 95.238\%$

**Character x** has the worst performance -  $17/35 = 0.48$

Accuracy Rate =  $0.48 \times 100 = 48\%$

2.d) d. What was the total accuracy on the test and train sets?

Ans:

```
> accuracy_test <- sum(diag(conf_matrix_test))/sum(conf_matrix_test)
> print(accuracy_test)
```

**Output: 0.76 or 76%**

```
> accuracy_test <- sum(diag(conf_matrix_test))/sum(conf_matrix_test)
> print(accuracy_test)
[1] 0.76
> |
```

```
> test <- az_char[-indices,] ##the 20% data to be assigned to test_data
```



## Accuracy for Training data

```
> training <- az_char[indices] ##80% sampled data
> predict_lda_training <- predict(object = lda_model, newdata = training)
> conf_matrix_training <- table(training$char, predict_lda_training$class)
> print(conf_matrix_training)
```

```
Console C:/Users/dubey/OneDrive/Quarter3/Pattern/assign2/assign2/
> test <- az_char[-indices,]
> training <- az_char[indices,]
> predict_lda_training <- predict(object = lda_model, newdata = training)
> conf_matrix_training <- table(training$char, predict_lda_training$class)
> print(conf_matrix_training)
```

	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z
a	85	0	2	15	2	2	1	1	0	0	1	2	0	0	0	0	1	0	0	1	1	0	1	7	2	16
b	0	118	0	0	5	3	0	5	0	0	0	2	0	0	0	0	0	0	1	1	0	0	0	0	1	0
c	1	0	144	0	6	0	0	3	0	0	0	2	1	0	0	0	0	0	0	1	0	0	1	0	0	0
d	12	1	3	100	1	0	0	1	4	0	2	2	2	1	2	0	0	0	0	0	1	0	0	23	0	0
e	0	0	3	0	151	0	0	1	0	0	0	0	0	0	0	2	0	0	0	1	0	0	0	1	1	0
f	1	0	0	0	0	93	1	0	0	0	1	6	0	0	0	14	0	2	0	10	0	0	0	5	2	0
g	1	0	0	0	0	12	91	0	0	0	0	0	0	0	0	0	5	0	21	5	1	0	0	2	8	1
h	4	4	0	0	2	0	0	111	0	0	5	8	1	6	0	0	0	0	0	0	3	0	1	0	0	1
i	0	0	0	0	0	3	0	0	114	14	0	8	0	0	0	2	0	1	0	12	0	0	0	0	0	0
j	0	2	0	0	0	0	0	0	18	117	0	1	0	0	1	0	0	1	4	5	0	0	0	2	5	0
k	0	1	0	0	1	1	0	19	0	0	86	4	6	2	0	0	0	0	0	3	5	1	1	1	2	0
l	0	0	1	6	10	4	0	12	0	0	0	142	0	0	0	0	0	0	0	0	0	0	0	3	0	0
m	1	0	0	1	0	0	0	0	0	0	1	0	136	5	0	0	1	0	0	0	6	0	2	0	0	0
n	3	0	0	2	0	0	0	5	0	0	1	0	2	98	0	0	1	1	0	0	14	0	3	0	0	1
o	1	0	1	4	0	4	4	0	0	0	0	0	0	0	152	1	0	2	2	1	0	2	2	0	0	0
p	0	0	0	0	0	1	0	0	0	0	0	1	0	1	0	151	1	7	1	1	0	0	0	0	1	0
q	1	0	0	0	4	6	16	0	0	0	0	3	0	0	0	125	0	1	1	0	0	0	2	8	0	0
r	0	0	0	0	0	1	0	1	2	0	0	0	2	1	0	4	3	146	0	2	1	5	0	0	0	1
s	0	0	0	0	0	0	6	0	0	0	0	2	1	2	0	0	0	1	147	2	0	0	0	0	3	0
t	0	2	0	0	0	9	0	0	4	0	0	2	0	0	3	0	0	5	0	122	0	0	1	5	0	0
u	8	0	0	1	0	0	0	9	0	0	3	0	2	4	0	1	0	0	0	0	123	24	3	0	0	0
v	0	0	0	0	0	0	0	0	0	1	1	0	0	0	2	0	0	4	0	0	7	149	1	0	0	0
w	0	0	0	0	0	0	0	0	0	0	1	0	1	0	0	0	0	1	0	0	0	4	145	0	0	0
x	5	3	7	9	3	0	2	0	0	0	3	1	0	2	3	0	0	3	1	3	2	18	0	72	15	0
y	0	0	0	0	0	3	2	0	0	0	0	4	0	0	0	0	0	0	3	5	0	1	0	0	138	0
z	0	2	0	0	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	3	0	0	0	2	1	110

```
> accuracy_training <- sum(diag(conf_matrix_training))/sum(conf_matrix_training)
> print(accuracy_training)
```

```
> accuracy_training <- sum(diag(conf_matrix_training))/sum(conf_matrix_training)
> print(accuracy_training)
[1] 0.7915
> |
```

The accuracy on training is 79.2% which is better than the test data.

3. Logistic Regression. The file "credit\_data.txt" contains information about the financial characteristics of 885 firms, which applied for a bank loan. Use the sample command to randomly select 80% of the data for training. Use the table command to show the number of cases per class in the training and test data.

```
> credit_data <- read.table("C:/Users/dubey/OneDrive/Quarter3/Pattern/assign2/credit_data.txt", header = T) ## file read
> sampling <- sample(1:nrow(credit_data), size = (0.8*nrow(credit_data))) ##80% data sampled used for training
> train_data <- credit_data[sampling,]
```

```
>test_data<-credit_data[-sampling,]# testing data
```

Train\_data with 708 observations of 15 variables

Id	Fail	Leverage	CumulProfit	Liquid	OverDueDebt	WorkCap	OperProfit	ShortDebt	GuarDebt	StateLag	FiscalLag	InFinan	Links
2	2	0	1	0.08815	0.02957	0	0.03307	0.13316	0.01061	0.04186	136	35	0.09217
10	10	0	1	0.02211	0.02029	0	0.09951	0.05124	0.02020	0.00000	206	160	0.04400
11	11	0	0	0.05116	0.00870	0	0.54226	0.07555	0.34457	0.00000	184	177	0.08328
14	14	1	0	0.01083	0.00294	1	0.01507	0.10429	0.15281	0.83896	366	311	0.20141
17	17	0	1	0.15301	0.30415	0	0.19971	0.14883	0.00000	0.00000	172	151	0.16591
19	19	0	1	0.09558	0.25440	0	-0.13621	-0.01643	0.00000	0.00000	247	100	-0.03902
23	23	0	1	0.31044	0.00982	0	0.41486	0.08323	0.06233	0.00000	176	109	0.00895
27	27	1	0	-0.07613	0.00320	1	0.57717	0.06571	0.42575	0.00000	179	72	0.06632
28	28	0	1	0.46627	0.05241	0	0.64956	0.11040	0.17110	0.00000	149	131	0.07357
29	29	0	1	0.21735	0.11486	0	0.44993	0.12197	0.38506	0.65291	243	174	0.15048
31	31	0	0	0.09442	0.01080	0	-0.00286	0.05810	0.00000	0.35419	333	148	-0.01063
36	36	0	1	0.36821	0.30347	0	0.05772	0.22207	0.00000	0.17524	167	148	0.03095
52	52	1	1	0.04537	0.02590	1	0.17760	-0.05643	0.12111	0.00000	289	116	0.23265
55	55	0	1	0.31327	0.24581	0	0.34631	-0.00875	0.00000	0.00000	125	98	0.07935
58	58	1	0	-0.12942	0.00043	0	0.30786	-0.01843	0.17119	0.33301	158	130	-0.04411
70	70	0	0	0.07263	0.04743	0	0.61638	0.00277	0.59292	0.00000	163	132	0.03280
80	80	0	1	0.51462	0.14408	0	0.43635	0.33254	0.00000	0.00000	175	151	0.36931
93	93	1	0	-0.00296	0.04848	1	0.07505	-0.10625	0.07249	0.52523	186	156	0.13988
105	105	0	0	0.14052	0.04983	0	0.10704	-0.00464	0.56701	0.00000	284	59	-0.49910
109	109	0	1	0.18332	0.00067	0	0.35893	0.04088	0.00074	0.61804	199	143	0.06659
112	112	0	0	0.21065	0.00239	0	0.11854	0.01996	0.18929	0.44041	170	136	0.05777
116	116	0	1	0.20049	0.02141	0	0.35451	0.19577	0.17479	0.00000	232	128	0.21197

Test\_data with 177 observations out of 15 variables.

Id	Fail	Leverage	CumulProfit	Liquid	OverDueDebt	WorkCap	OperProfit	ShortDebt	GuarDebt	StateLag	FiscalLag	InFinan	Links
2	2	0	1	0.08815	0.02957	0	0.03307	0.13316	0.01061	0.04186	136	35	0.09217
10	10	0	1	0.02211	0.02029	0	0.09951	0.05124	0.02020	0.00000	206	160	0.04400
11	11	0	0	0.05116	0.00870	0	0.54226	0.07555	0.34457	0.00000	184	177	0.08328
14	14	1	0	0.01083	0.00294	1	0.01507	0.10429	0.15281	0.83896	366	311	0.20141
17	17	0	1	0.15301	0.30415	0	0.19971	0.14883	0.00000	0.00000	172	151	0.16591
19	19	0	1	0.09558	0.25440	0	-0.13621	-0.01643	0.00000	0.00000	247	100	-0.03902
23	23	0	1	0.31044	0.00982	0	0.41486	0.08323	0.06233	0.00000	176	109	0.00895
27	27	1	0	-0.07613	0.00320	1	0.57717	0.06571	0.42575	0.00000	179	72	0.06632
28	28	0	1	0.46627	0.05241	0	0.64956	0.11040	0.17110	0.00000	149	131	0.07357
29	29	0	1	0.21735	0.11486	0	0.44993	0.12197	0.38506	0.65291	243	174	0.15048
31	31	0	0	0.09442	0.01080	0	-0.00286	0.05810	0.00000	0.35419	333	148	-0.01063
36	36	0	1	0.36821	0.30347	0	0.05772	0.22207	0.00000	0.17524	167	148	0.03095
52	52	1	1	0.04537	0.02590	1	0.17760	-0.05643	0.12111	0.00000	289	116	0.23265
55	55	0	1	0.31327	0.24581	0	0.34631	-0.00875	0.00000	0.00000	125	98	0.07935
58	58	1	0	-0.12942	0.00043	0	0.30786	-0.01843	0.17119	0.33301	158	130	-0.04411
70	70	0	0	0.07263	0.04743	0	0.61638	0.00277	0.59292	0.00000	163	132	0.03280
80	80	0	1	0.51462	0.14408	0	0.43635	0.33254	0.00000	0.00000	175	151	0.36931
93	93	1	0	-0.00296	0.04848	1	0.07505	-0.10625	0.07249	0.52523	186	156	0.13988
105	105	0	0	0.14052	0.04983	0	0.10704	-0.00464	0.56701	0.00000	284	59	-0.49910
109	109	0	1	0.18332	0.00067	0	0.35893	0.04088	0.00074	0.61804	199	143	0.06659
112	112	0	0	0.21065	0.00239	0	0.11854	0.01996	0.18929	0.44041	170	136	0.05777
116	116	0	1	0.20049	0.02141	0	0.35451	0.19577	0.17479	0.00000	232	128	0.21197

```
#case by class(fail) for test_data
```

```
>case_of_test_data<- table(test_data$Fail)
```

```
>print(case_of_test_data)
```

```
> case_of_test_data<- table(test_data$Fail)
> print(case_of_test_data)

 0  1
137 40
> |
```

```
>case_of_train_data<-table(train_data$Fail)
>print(case_of_train_data)
> case_of_train_data<-table(train_data$Fail)
> print(case_of_train_data)

  0    1
556 152
```

(a) Use the glm (with family=binomial) command to fit a logistic regression to predict which firms will go bankrupt. Report the table of coefficients from R with their p-values. What are the 4 most important predictor variables?

Ans:

```
>lr_training_data <- glm(formula = Fail ~.,family=binomial,data=train_data)
>model_lr_train <- predict(object = lr_training_data, newdata = train_data)
>print(model_lr_train)
> lr_training_data <- glm(formula = Fail ~.,family=binomial,data=train_data)
> model_lr_train <- predict(object = lr_training_data, newdata = train_data)
> print(model_lr_train)

      670      333      94      544      745      59      132      349
-2.922805891 -0.452860410 -8.539188687 -1.636920483 -2.369191867 -7.998400607 -1.312800569 -2.321437024
      481      725      836      282      715      357      601      194
-0.473301223 -1.137626694 -0.254599338 -1.743480916 -0.691742244 -3.054726285 -3.624532485 -2.629521205
      698      472      42      341      20      791      191      846
-0.399025910 -1.259692210 -2.328017678 -2.028641620 -3.527489922 -0.002648529 -1.782767419 1.699665164
      92      169      540      721      142      652      139      879
-5.503250004 -1.978432102 -0.712317314 -0.216700124 -0.756353624 1.557707237 -4.883294022 -2.297018344
      290      841      564      589      260      692      374      30
-3.460427708 -0.280154289 -5.475031094 -0.842821252 -3.114134871 -1.115859304 -2.357130066 -3.975102426
      351      241      547      876      656      880      815      368
-2.624159551 -4.379615055 -0.233960369 1.499565643 -1.313999530 -0.978209012 0.724704158 -1.027504270
      464      56      735      218      82      789      367      298
-3.478670510 -2.212055127 -0.386390178 -1.014074357 -1.907073292 -0.497409311 -6.675074376 -2.923802511
      53      198      873      22      304      38      864      730
-6.321119070 -0.865667880 0.748184359 -5.197412622 -3.161234041 -4.602333327 -1.316738896 -1.406031701
      208      505      635      272      369      373      840      613
-5.193817338 -5.519104274 -0.724311930 -3.827861436 -6.937445042 -1.564774197 1.727144330 -1.411240580
      15      197      311      623      617      353      431      365
-2.137768356 -3.976589473 -3.928062302 -2.397510144 -6.767704764 -3.795595663 -14.125355747 -1.078139412
      339      133      214      667      772      281      583      343
-5.846193101 -5.385935224 -3.201359598 -2.709547719 4.636070703 -2.962963366 -3.046965332 -0.888176555
      600      669      548      334      13      278      496      155
-8.267387205 -2.276894451 -0.586249995 -1.972214680 -0.663559732 -2.963904561 0.044811377 -2.961376200
      1      270      602      517      215      428      662      165
-0.174373060 -1.055104333 0.561108813 0.333045635 0.314331380 1.114333131 0.865330073 0.811135533
```

```
>coef(summary(lr_training_data))
```



```
> coef(summary(lr_training_data))
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-4.557556878	0.9031007756	-5.0465651	4.498236e-07
Id	0.002148536	0.0004681292	4.5896219	4.440495e-06
Leverage	-0.493995294	0.2466330535	-2.0029566	4.518195e-02
CumulProfit	-1.692239179	0.6903486445	-2.4512820	1.423484e-02
Liquid	-16.579712590	3.5296816244	-4.6972261	2.637185e-06
OverDueDebt	1.502811194	0.3167887305	4.7438910	2.096517e-06
WorkCap	-1.038433574	0.5859744658	-1.7721482	7.636997e-02
OperProfit	-0.285993727	0.5465026472	-0.5233163	6.007542e-01
ShortDebt	2.080630290	0.6522781399	3.1897900	1.423762e-03
GuarDebt	-1.043623812	0.4813653221	-2.1680494	3.015493e-02
StateLag	0.008573054	0.0022082772	3.8822366	1.035001e-04
FiscalLag	-0.003708271	0.0031337794	-1.1833223	2.366815e-01
InFinan	-0.378553989	0.5130731751	-0.7378168	4.606258e-01
Links	6.036591473	2.6616091602	2.2680233	2.332779e-02
CapStruct	1.894084645	0.7488028346	2.5294838	1.142304e-02

```
>
```

>tail(sort(abs(coef(summary(model))[,1])), 4) ## intercept is included in the 4 important values, **taking the absolute value**

```
> tail(sort(abs(coef(summary(model))[,1])), 4)
```

ShortDebt	(Intercept)	Links	Liquid
2.080630	4.557557	6.036591	16.579713

>tail(sort(abs(coef(summary(model))[,1])), 5) ##taking the absolute values of the estimate

**(Note: Here we are taking the 5 most important values because intercept is one of them)**

```
> tail(sort(abs(coef(summary(model))[,1])), 5)
```

CapStruct	ShortDebt	(Intercept)	Links	Liquid
1.894085	2.080630	4.557557	6.036591	16.579713

b)Do their signs appear to be what you'd expect?

**Ans: Yes, according to us these variables play an important role in predicting whether a company would earn profits or go bankrupt. Clearly, parameter Liquid plays the most important factor in predicting about the bankruptcy.**

(c) Suppose that we predict a firm will go bankrupt if the predicted probability  $P(Y = 1 | X = x)$  of bankruptcy is 0.5 or greater. Find the confusion matrix for such predictions on the test data.

Ans:

```
>new_predict_test <- rep("0", nrow(test_data))
>new_predict_test[predict_model_test >= 0.5] <- "1"
>print(new_predict_test)
>cm <- table(test_data$Fail, new_predict_test)
>print(cm)
>accuracy <- sum(diag(cm))/sum(cm)
>print(accuracy)
```

```

> new_predict_test <- rep("0", nrow(test_data))
> new_predict_test[predict_model_test >= 0.5] <- "1"
> print(new_predict_test)
[1] "0" "0" "0" "0" "0" "0" "0" "1" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
[28] "0" "0" "0" "0" "0" "1" "0" "0" "0" "0" "0" "0" "0" "1" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
[55] "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "1" "0" "0" "0"
[82] "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "1" "0" "0" "0" "0" "1"
[109] "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "1" "0" "0" "0" "0" "0" "1" "0" "0"
[136] "0" "0" "1" "0" "0" "0" "0" "0" "0" "1" "0" "0" "0" "0" "0" "0" "0" "0" "1" "1" "1" "1" "0" "1" "1" "0"
[163] "0" "0" "1" "1" "0" "1" "1" "1" "1" "0" "1" "1" "1" "1" "0"
> cm <- table(test_data$Fail, new_predict_test)
> accuracy <- sum(diag(cm))/sum(cm)
> print(cm)
      new_predict_test
      0      1
0 129      8
1   20     20
> print(accuracy)
[1] 0.8418079

```

**Accuracy = 84.18**

4. Regularized Logistic Regression. The R package glmnet fits penalized logistic regression models using the Lasso penalty. We want to compare the regularized vs. the unregularized fit to the credit data.

(a) Use the cv.glmnet (with family=binomial) command to fit a regularized logistic regression to the same training data used in 3a (you may need a cast from data.frame to matrix and map y from 0/1 to -1/1). Plot the cross-validation curve. Explain the plot.

Ans:

```

> library(glmnet)
> train_data_matrix <- data.matrix(subset(train_data, select = -c(Fail)))
> train_data_y <- train_data[2] ## extracting the ~Fail column
> train_data_y[train_data_y == 0] <- "-1" ##changing the values of Y_matrix, 0 to -1
> train_data_y <- data.matrix(train_data_y)
> cv_out <- cv.glmnet(x = train_data_matrix, y = train_data_y, alpha = 1, family = "binomial")

```

```


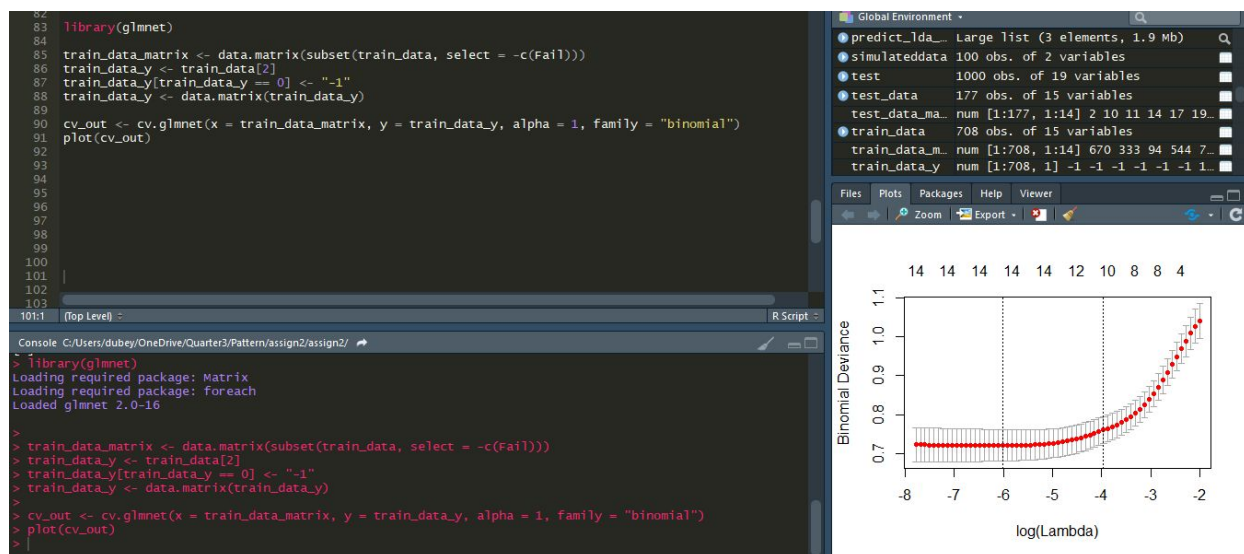
> train_data_y <- data.matrix(train_data_y)
> #matrix_train_data <- model.matrix(Fail~.,train_data)
> train_data_matrix <- data.matrix(subset(train_data, select = -c(Fail)))
> train_data_y <- train_data[2]
> train_data_y[train_data_y == 0] <- "-1"
> train_data_y <- data.matrix(train_data_y)
> cv_out <- cv.glmnet(x = train_data_matrix, y = train_data_y, alpha = 1, family = "binomial")
>

```

```

> plot(cv_out)

```



Explanation - `cv.glmnet()` uses cross-validation to determine how each model used inside the function behaves in the dataset which can be visualised from the above plot. The above plot shows that the  $\log(\lambda)$  fits best to the values from -6.5 to -5.8. The lowest point in the curve indicates the optimal lambda: the log value of lambda that best minimised the error in cross-validation.

(b) The object returned by `cv.glmnet()` contains the value of the best lambda. Pass this value of lambda to the `coef()` function to retrieve the corresponding coefficient vector. Print the coefficients. Compare to your answer in 3a.

Ans:



```
>coef(cv_out)
```

```
> coef(cv_out)
15 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept) -3.990932086
Id           0.001712782
Leverage     -0.378537917
CumulProfit -0.888536430
Liquid       -4.747562235
OverDueDebt  1.001479324
WorkCap      .
OperProfit   .
ShortDebt    0.955522736
GuarDebt     .
StateLag     0.004184330
FiscalLag    .
InFinan      -0.045374038
Links        0.878463181
CapStruct    1.363400681
```

Below is the coefficient of 3(a). Here we get coefficient value in Workcap, OperProfit, GuarDebt, FiscalLag unlike above.

```
> coef(model)
      (Intercept)      Id      Leverage      CumulProfit      Liquid      overDueDebt      workCap
-4.557556878    0.002148536 -0.493995294 -1.692239179 -16.579712590  1.502811194 -1.038433574
      OperProfit      ShortDebt      GuarDebt      StateLag      FiscalLag      InFinan      Links
-0.285993727    2.080630290 -1.043623812  0.008573054  -0.003708271 -0.378553989  6.036591473
      CapStruct
1.894084645
```

`>cv_out[["glmnet.fit"]][["beta"]]`## extracts the all the coefficients of the models predicted by Lasso regularization giving us beta values of 64 tested models.

Console C:\Users\adobej\OneDrive\Quarter3\Future\Assign2\Assign2

```
> cv_out[["glmnet.fit"]][["beta"]]
14 x 64 sparse Matrix of class "dgcMatrix"
[[ suppressing 64 column names 's0', 's1', 's2' ... ]]

Id      .      .      0.0001779567  0.0003772601  0.0005547713  0.0007026127  0.0008342062
Leverage .      .      .      .      .      .      .
CumulProfit . -0.1476268 -0.2298587436 -0.3008665159 -0.3682696335 -0.4062472025 -0.4354373352
Liquid .      .      .      .      .      .      .
OverDueDebt . .      .      .      .      .      .
WorkCap .      .      .      .      .      .      .
OperProfit . .      .      .      .      .      .
ShortDebt . .      .      .      .      .      .
GuarDebt . .      .      .      .      .      .
StateLag . .      .      .      .      .      .
FiscalLag . .      .      .      .      .      .
InFinan . .      .      .      .      .      .
Links .      .      .      .      .      .      .
CapStruct . 0.1109991 0.2340852002 0.3464057248 0.4485375463 0.5516182243 0.6542732700

Id      0.0009401143 0.0010344917 0.001121520 0.001202126 0.001277042 0.001346843
Leverage -0.1456711275 -0.1766877216 -0.204766055 -0.230248821 -0.253427518 -0.274536046
CumulProfit -0.4634928398 -0.4943855986 -0.527215128 -0.561521934 -0.596904425 -0.633092796
Liquid -0.2104573204 -0.4876629112 -0.781412800 -1.093381795 -1.424816456 -1.776570507
OverDueDebt 0.3893332625 0.4656860466 0.535032577 0.598544748 0.657083526 0.711299493
WorkCap .      .      .      .      .      .
OperProfit .      .      .      .      .      .
ShortDebt 0.0969921627 0.2098277724 0.312082464 0.404985452 0.489519111 0.566509828
GuarDebt .      .      .      .      .      .
StateLag 0.0004440165 0.0008905403 0.001301539 0.001682793 0.002038644 0.002372365
FiscalLag .      .      .      .      .      .
InFinan .      .      .      .      .      .
Links .      .      .      .      .      .
CapStruct 0.7244925256 0.7853189308 0.844875351 0.903089356 0.959927465 1.015340934

Id      0.001412003 0.001472908 0.001529887 0.001583240 0.001632396 0.001674602
Leverage -0.293791158 -0.311360219 -0.327392207 -0.342041735 -0.355689053 -0.368420113
CumulProfit -0.669823795 -0.706969295 -0.744424795 -0.782032121 -0.819603877 -0.855334683
Liquid -2.148975809 -2.541955716 -2.954918332 -3.386659241 -3.834632043 -4.289641022
OverDueDebt 0.761692552 0.808663964 0.852541768 0.893597942 0.931929399 0.967604533
```

```
Liquid -0.2104573204 -0.4876629112 -0.781412800 -1.093381795 -1.424816456 -1.776570507
OverDueDebt 0.3893332625 0.4656860466 0.535032577 0.598544748 0.657083526 0.711299493
WorkCap .      .      .      .      .      .
OperProfit .      .      .      .      .      .
ShortDebt 0.0969921627 0.2098277724 0.312082464 0.404985452 0.489519111 0.566509828
GuarDebt .      .      .      .      .      .
StateLag 0.0004440165 0.0008905403 0.001301539 0.001682793 0.002038644 0.002372365
FiscalLag .      .      .      .      .      .
InFinan .      .      .      .      .      .
Links .      .      .      .      .      .
CapStruct 0.7244925256 0.7853189308 0.844875351 0.903089356 0.959927465 1.015340934

Id      0.001412003 0.001472908 0.001529887 0.001583240 0.001632396 0.001674602
Leverage -0.293791158 -0.311360219 -0.327392207 -0.342041735 -0.355689053 -0.368420113
CumulProfit -0.669823795 -0.706969295 -0.744424795 -0.782032121 -0.819603877 -0.855334683
Liquid -2.148975809 -2.541955716 -2.954918332 -3.386659241 -3.834632043 -4.289641022
OverDueDebt 0.761692552 0.808663964 0.852541768 0.893597942 0.931929399 0.967604533
WorkCap .      .      .      .      .      .
OperProfit .      .      .      .      .      .
ShortDebt 0.636655519 0.700588209 0.758877019 0.812029090 0.861094021 0.909182223
GuarDebt .      .      .      .      .      .
StateLag 0.002686429 0.002982724 0.003262706 0.003527511 0.003773708 0.003987319
FiscalLag .      .      .      .      .      .
InFinan .      .      .      .      .      .
Links .      .      .      .      .      .
CapStruct 1.069343883 1.121904840 1.173012334 1.222711202 1.271160080 1.318351866

Id      0.001712782 0.001748641 0.001782339 0.001811533 0.001836528 0.001860163
Leverage -0.378537917 -0.387876933 -0.396490352 -0.404358939 -0.411480118 -0.418073540
CumulProfit -0.888536430 -0.922708823 -0.957585605 -0.988729917 -1.023074731 -1.058072923
Liquid -4.747562235 -5.215012902 -5.689055609 -6.187862166 -6.710262141 -7.227054070
OverDueDebt 1.001479324 1.033589353 1.064004079 1.092315910 1.118530254 1.143577888
WorkCap .      .      .      .      .      .
OperProfit .      .      .      .      .      .
ShortDebt 0.955522736 0.997928854 1.036816106 1.081546238 1.152105805 1.218039999
GuarDebt .      .      .      .      .      .
StateLag 0.004184330 0.004371057 0.004547889 0.004716926 0.004884375 0.005043971
FiscalLag .      .      .      .      .      .
InFinan -0.045374038 -0.076508527 -0.105283152 -0.136071002 -0.149078720 -0.160928307
Links .      .      .      .      .      .
CapStruct 0.970463183 1.038524003 1.106165270 1.173030636 1.230601316 1.288433781
```



CumulProfit	-0.888536430	-0.922708823	-0.957585605	-0.988729917	-1.023074731	-1.058072923
Liquid	-4.747562235	-5.215012902	-5.689055609	-6.187862166	-6.710262141	-7.227054070
OverDueDebt	1.001479324	1.033589353	1.064004079	1.092315910	1.118530254	1.143577888
workCap	.	.	.	-0.012447559	-0.083863274	-0.150928243
OperProfit	.	.	.	.	.	.
ShortDebt	0.955522736	0.997928854	1.036816106	1.081546238	1.152105805	1.218039999
GuarDebt	.	.	.	-0.058188653	-0.125109988	-0.187754039
StateLag	0.004184330	0.004371057	0.004547889	0.004716926	0.004884375	0.005043971
FiscalLag	.	.	.	.	.	.
InFinan	-0.045374038	-0.076508527	-0.105283152	-0.136071002	-0.149078720	-0.160928307
Links	0.878463181	1.228574903	1.551657794	1.870030635	2.172060125	2.455422781
CapStruct	1.363400681	1.407297522	1.450023823	1.489911603	1.504854656	1.520219928
Id	0.001882541	0.001903922	0.001923972	0.001942397	0.001959350	0.001975409
Leverage	-0.424198392	-0.430205074	-0.435481574	-0.439698034	-0.443025971	-0.446131060
CumulProfit	-1.093308059	-1.127519069	-1.162362344	-1.194800176	-1.225538268	-1.255300506
Liquid	-7.735682762	-8.222051111	-8.707992207	-9.178238268	-9.631993612	-10.067318809
OverDueDebt	1.167459281	1.189984410	1.211560631	1.231866065	1.250974489	1.268991196
workCap	-0.213758299	-0.271779771	-0.326816406	-0.378061253	-0.425950939	-0.470324744
OperProfit	.	.	.	-0.009705785	-0.026083959	-0.041677595
ShortDebt	1.279613785	1.336337360	1.390081241	1.439969248	1.486466575	1.529684969
GuarDebt	-0.246375379	-0.301062922	-0.352313510	-0.400615254	-0.446018533	-0.488338262
StateLag	0.005195811	0.005338593	0.005475382	0.005602722	0.005721415	0.005833134
FiscalLag	.	.	.	.	.	.
InFinan	-0.171775378	-0.181835012	-0.190994713	-0.200522927	-0.210131357	-0.219137023
Links	2.721400533	2.970464486	3.204724873	3.423324956	3.627258544	3.818022952
CapStruct	1.535952614	1.551697139	1.567721268	1.584159215	1.600603869	1.616978124
Id	0.001990566	0.002002124	2.012242e-03	2.021813e-03	2.030854e-03	0.002039378
Leverage	-0.448982627	-0.452325007	-4.556394e-01	-4.587212e-01	-4.615827e-01	-0.464236367
CumulProfit	-1.284190171	-1.311620840	-1.337798e+00	-1.362790e+00	-1.386545e+00	-1.409034710
Liquid	-10.484284480	-10.893896408	-1.128789e+01	-1.166227e+01	-1.201672e+01	-12.351317050
OverDueDebt	1.285962595	1.301445700	1.315929e+00	1.329588e+00	1.342437e+00	1.354498495
workCap	-0.511645699	-0.552198449	-5.905323e-01	-6.261246e-01	-6.591264e-01	-0.689698637
OperProfit	-0.056512881	-0.070762718	-8.437064e-02	-9.731798e-02	-1.096246e-01	-0.121309599
ShortDebt	1.569982821	1.607826216	1.643341e+00	1.676559e+00	1.707576e+00	1.736500222
GuarDebt	-0.527757659	-0.566744530	-6.037632e-01	-6.382302e-01	-6.702741e-01	-0.700030994
StateLag	0.005938189	0.006112646	6.298826e-03	6.473076e-03	6.635693e-03	0.006787232
FiscalLag	.	-0.000223023	-4.955668e-04	-7.490701e-04	-9.844543e-04	-0.001202783
FiscalLag	.	-0.000223023	-4.955668e-04	-7.490701e-04	-9.844543e-04	-0.001202783
InFinan	-0.227515301	-0.238401116	-2.494241e-01	-2.596768e-01	-2.692025e-01	-0.278040422
Links	3.996330965	4.154230674	4.298930e+00	4.433724e+00	4.559170e+00	4.675800099
CapStruct	1.632964555	1.648591985	1.663886e+00	1.678736e+00	1.693057e+00	1.706975390
Id	0.002047402	0.002054939	0.002062007	0.002068623	0.002074805	0.002080571
Leverage	-0.466694057	-0.468967202	-0.471066847	-0.473003702	-0.474788143	-0.476430223
CumulProfit	-1.430249588	-1.450197216	-1.468898344	-1.486384378	-1.502695001	-1.517876116
Liquid	-12.666329073	-12.962156378	-13.239318167	-13.498427683	-13.740171381	-13.965290150
OverDueDebt	1.365797728	1.376362298	1.386221872	1.395407566	1.403951476	1.411886252
workCap	-0.717994684	-0.744160580	-0.768335347	-0.790651243	-0.811233962	-0.830202783
OperProfit	-0.132391303	-0.142887203	-0.152814493	-0.162190285	-0.171031805	-0.179356528
ShortDebt	1.763437898	1.788493390	1.811768625	1.833363672	1.853375943	1.871900245
GuarDebt	-0.276324681	-0.293206780	-0.307687606	-0.3208761067	-0.332974529	-0.344026683
StateLag	0.006928249	0.007059293	0.007180910	0.007293634	0.007397990	0.007494488
FiscalLag	-0.001405080	-0.001592332	-0.001765484	-0.001925444	-0.002073081	-0.002209223
InFinan	-0.286229199	-0.293806783	-0.300810282	-0.307275791	-0.313238237	-0.318731254
Links	4.784121106	4.884616465	4.977750598	5.063970072	5.143704371	5.217366159
CapStruct	1.719913983	1.732387681	1.744202676	1.755354734	1.765847680	1.775692002
Id	0.002090932	0.002090932	0.002095566	0.002099891	0.002103867	0.002107542
Leverage	-0.477939607	-0.479325593	-0.480597064	-0.481818784	-0.482884480	-0.483857771
CumulProfit	-1.531978069	-1.545054144	-1.557159305	-1.568217718	-1.578349091	-1.588080758
Liquid	-14.174562650	-14.368790749	-14.548786994	-14.712582877	-14.866555722	-15.008820093
OverDueDebt	1.419244723	1.426059554	1.432362970	1.438143802	1.443518132	1.448474708
workCap	-0.847670706	-0.863744587	-0.878525272	-0.891950778	-0.904428064	-0.915885528
OperProfit	-0.187182271	-0.194527233	-0.201409992	-0.207822011	-0.213836978	-0.219448474
ShortDebt	1.889028470	1.904849359	1.919448286	1.932761620	1.945160756	1.956577803
GuarDebt	-0.854822442	-0.870661828	-0.885239946	-0.898588123	-0.910909058	-0.922227010
StateLag	0.007583622	0.007665871	0.007741694	0.007811074	0.007875341	0.007934457
FiscalLag	-0.002334661	-0.002450143	-0.002556380	-0.002653787	-0.002743502	-0.002825880
InFinan	-0.323787079	-0.328436481	-0.332708704	-0.336646839	-0.340245467	-0.343544950
Links	5.285351245	5.348038405	5.405789164	5.458691385	5.507583289	5.552529735
CapStruct	1.784903564	1.793502461	1.801512001	1.808822843	1.815735150	1.822137948
Id	0.002110937	0.002114069	0.002116956	0.002119614	0.002122077	0.002124325
Leverage	-0.484747770	-0.485561235	-0.486304401	-0.486983064	-0.487650726	-0.488215737
CumulProfit	-1.596860419	-1.604939150	-1.6123666178	-1.619188471	-1.625340618	-1.631083684
Liquid	-15.140031090	-15.260921822	-15.372200500	-15.474543847	-15.565633704	-15.652002562



```

OperProfit -0.224674981 -0.229535939 -0.234050714 -0.238238457 -0.242086147 -0.245675355
ShortDebt 1.967078000 1.976727810 1.985590118 1.993724079 2.001041684 2.007882640
GuardDebt -0.932615205 -0.942144369 -0.950880789 -0.958886377 -0.966150934 -0.972863954
StateLag 0.007988780 0.008038663 0.008084442 0.008126428 0.008164395 0.008199659
FiscalLag -0.002901469 -0.002970790 -0.003034331 -0.003092545 -0.003145538 -0.003194336
InFinan -0.346568798 -0.349338601 -0.351874487 -0.354195196 -0.356330255 -0.358271127
Links 5.593812342 5.631704940 5.666465086 5.698334028 5.727244033 5.753991328
CapStruct 1.828065239 1.833541535 1.838594850 1.843252574 1.847371575 1.851317300

Id 0.002126389 0.002128283 0.002130021 0.002131623 0.002133084 0.002134421
Leverage -0.488728014 -0.489195000 -0.489620804 -0.490050950 -0.490406646 -0.490726137
CumulProfit -1.63356617 -1.641188508 -1.645611354 -1.649585406 -1.653285559 -1.656680335
Liquid -15.731457992 -15.804350363 -15.871172934 -15.929176797 -15.985140595 -16.036688643
OverDueDebt 1.473605586 1.476123145 1.478429684 1.480488363 1.482420262 1.484192987
WorkCap -0.973206484 -0.978889844 -0.984086589 -0.988690585 -0.993029141 -0.997005755
OperProfit -0.248994905 -0.252060036 -0.254887293 -0.257458465 -0.259855484 -0.262064335
ShortDebt 2.014160193 2.019910704 2.025175637 2.029848363 2.034252535 2.038294014
GuardDebt -0.979011759 -0.984636406 -0.989780423 -0.994403308 -0.998699668 -1.002632840
StateLag 0.008231986 0.008261582 0.008288664 0.008312819 0.008335460 0.008356212
FiscalLag -0.003239010 -0.003279877 -0.003317249 -0.003351021 -0.003382233 -0.003410788
InFinan -0.360044820 -0.361665715 -0.363146603 -0.364508383 -0.365743242 -0.366870270
Links 5.778497825 5.800927023 5.821446287 5.839872835 5.857018599 5.872721657
CapStruct 1.854954501 1.858296023 1.861363103 1.863956968 1.866530818 1.868909149

Id 0.002135645 0.002136765 0.002137790 0.002138700 0.002139500 0.002140200
Leverage -0.491016767 -0.491281542 -0.491556040 -0.491830538 -0.492105036 -0.492379534
CumulProfit -1.659785577 -1.662624145 -1.665188846 -1.667468344 -1.669463842 -1.671174340
Liquid -16.083896275 -16.127084888 -16.163188927 -16.192292975 -16.215407023 -16.232521071
OverDueDebt 1.485815083 1.487298418 1.488596307 1.489707196 1.490588085 1.491288974
WorkCap -1.000639654 -1.003958560 -1.006855475 -1.009302381 -1.011399289 -1.013146197
OperProfit -0.264095326 -0.265961121 -0.267640291 -0.269049160 -0.270198029 -0.271096898
ShortDebt 2.041990645 2.045369675 2.048317913 2.050849161 2.052944210 2.054681258
GuardDebt -1.006226855 -1.009509578 -1.012413985 -1.014929092 -1.017044199 -1.018759307
StateLag 0.008375185 0.008392521 0.008407626 0.008421631 0.008434536 0.008446341
FiscalLag -0.003436877 -0.003460705 -0.003481985 -0.003499713 -0.003513800 -0.003524247
InFinan -0.367899026 -0.368837954 -0.369697780 -0.370486508 -0.371125236 -0.371623964
Links 5.887074135 5.900186159 5.911781901 5.921977649 5.930874397 5.938581145
CapStruct 1.871089563 1.873085859 1.874634974 1.875849081 1.876730189 1.877385307

```

> which(cv\_out\$lambda == cv\_out\$lambda.min) ## to find out the minimum value of the lambda

The minimum value of the lambda, because the header row and column contains 0, we would extract the 43rd value from the Lasso model. Because whenever we run the program the value of the Lambda minimum changes therefore we put this in a variable called min\_lambda index and the value we need will be min\_lambda\_index+1

```

> which(cv_out$lambda == cv_out$lambda.min)
[1] 42

```

This is the Best Value of the lamda

```

> cv_out$lambda.min
[1] 0.002956224

```

0.002956

>min\_lambda\_index <- which(cv\_out\$lambda == cv\_out\$lambda.min)##extracting the coefficients of the 43rd row

>cv\_out\$glmnet.fit\$beta[,min\_lambda\_index+1]

>cv\_out\$glmnet.fit\$a0[min\_lambda\_index=1]

```

> min_lambda_index <- which(cv_out$lambda == cv_out$lambda.min)
> cv_out$glmnet.fit$beta[,min_lambda_index+1]
      Id      Leverage      CumulProfit      Liquid      OverDueDebt      WorkCap      OperProfit      ShortDebt      GuardDebt
0.002080571 -0.476430223 -1.517876116 -13.965290150 1.411886252 -0.830202783 -0.179356528 1.871900245 -0.837626685
      StateLag      FiscalLag      InFinan      Links      CapStruct
0.007494488 -0.002209223 -0.318731254 5.217366159 1.775692002
> cv_out$glmnet.fit$a0[min_lambda_index=1]
      s0
-1.296888
>

```

On extracting the Estimate values of the data we find that the

- Liquid(-13.96~14)
- Links(5.21)
- ShortDebt(1.87)
- CapStruct(1.77)

are the 4 most important predictor values. These values appear a to be approximately similar with the Estimate values in 3(a).

4.(c) Use the predict function with the same value of lambda to predict on the test data. Show the confusion matrix. Compare the accuracy with 3c.

Ans:

```
> new_cv_predict_class <- cv_out_predict_test >= cutoff #cutoff=0.5
> new_cv_predict_class <- as.numeric(new_cv_predict_class)
> cm_cv_test <- table(test_data$Fail, new_cv_predict_class)
> accuracy_cv <- sum(diag(cm_cv_test))/sum(cm_cv_test)
> print(accuracy_cv)
```

```
> new_cv_predict_class <- cv_out_predict_test >= cutoff
> new_cv_predict_class <- as.numeric(new_cv_predict_class)
> cm_cv_test <- table(test_data$Fail, new_cv_predict_class)
> accuracy_cv <- sum(diag(cm_cv_test))/sum(cm_cv_test)
> print(accuracy_cv)
[1] 0.8474576
```

The accuracy of the cross validated data is approximately the same as is in 3(c) which is 84.74.

6. We will now perform cross-validation on a simulated data set.

(a) Generate a simulated data set as follows:

```
> set.seed(1)
> y=rnorm(100)
> x=rnorm(100)
> y=x-2* x^2+ rnorm(100)
```


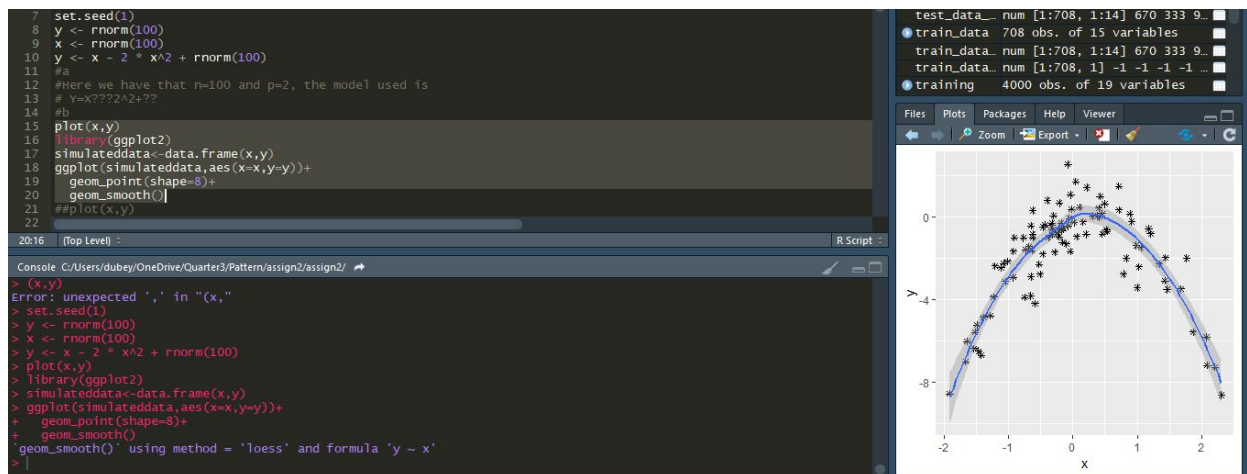
In this data set, what is n and what is p? Write out the model used to generate the data in equation form.

Ans:  $n = 100$ ,  $p = 2$ . The model used here is  $Y = X - 2 * x^2$ .

(b) Create a scatterplot of X against Y . Comment on what you find.

Ans:

```
> plot(x,y)
> library(ggplot2)
> simulateddata<-data.frame(x,y)
  ggplot(simulateddata,aes(x=x,y=y))+
  geom_point(shape=8)+
  geom_smooth()
```



**Plot explanation** - The plot here gives a parabolic curve or quadratic relationship between x and y.

(c) Set a random seed, and then compute the LOOCV errors that result from fitting the following four models using least squares:

i).  $Y = \beta_0 + \beta_1 X +$

Ans:

```

>library(boot)
>set.seed(1)
>Data <- data.frame(x, y)
>fit.glm.1 <- glm(y ~ x)
>cv.glm(Data, fit.glm.1)$delta[1]

```



```
> library(boot)
> set.seed(1)
> Data <- data.frame(x, y)
> fit.glm.1 <- glm(y ~ x)
> cv.glm(Data, fit.glm.1)$delta[1]
[1] 5.890979
```

ii.  $Y = \beta_0 + \beta_1 X + \beta_2 X^2 +$

Ans:

```
> fit.glm.2 <- glm(y ~ poly(x, 2))
> cv.glm(Data, fit.glm.2)$delta[1]
```

```
> fit.glm.2 <- glm(y ~ poly(x, 2))
> cv.glm(Data, fit.glm.2)$delta[1]
[1] 1.086596
```

iii.  $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 +$

Ans:

```
> fit.glm.3 <- glm(y ~ poly(x, 3))
> cv.glm(Data, fit.glm.3)$delta[1]
```

```
> fit.glm.3 <- glm(y ~ poly(x, 3))
> cv.glm(Data, fit.glm.3)$delta[1]
[1] 1.102585
```

iv.  $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \beta_4 X^4 +$ . Note you may find it helpful to use the data.frame() function to create a single data set containing both X and Y .

Ans:

```
> fit.glm.4 <- glm(y ~ poly(x, 4))
> cv.glm(Data, fit.glm.4)$delta[1]
```

```
> fit.glm.4 <- glm(y ~ poly(x, 4))
> cv.glm(Data, fit.glm.4)$delta[1]
[1] 1.114772
```

(d) Repeat (c) using another random seed, and report your results. Are your results the same as what you got in (c)? Why?

Ans:

```
> fit.glm.1 <- glm(y ~ x)
> cv.glm(Data, fit.glm.1)$delta[1]
```

```
> fit.glm.2 <- glm(y ~ poly(x, 2))
> cv.glm(Data, fit.glm.2)$delta[1]
```

```
> fit.glm.3 <- glm(y ~ poly(x, 3))
> cv.glm(Data, fit.glm.3)$delta[1]
```

```
> fit.glm.4 <- glm(y ~ poly(x, 4))
> cv.glm(Data, fit.glm.4)$delta[1]
```

```
> set.seed(10)
> fit.glm.1 <- glm(y ~ x)
> cv.glm(Data, fit.glm.1)$delta[1]
[1] 5.890979
> ## [1] 5.890979
> fit.glm.2 <- glm(y ~ poly(x, 2))
> cv.glm(Data, fit.glm.2)$delta[1]
[1] 1.086596
> ## [1] 1.086596
> fit.glm.3 <- glm(y ~ poly(x, 3))
> cv.glm(Data, fit.glm.3)$delta[1]
[1] 1.102585
> ## [1] 1.102585
> fit.glm.4 <- glm(y ~ poly(x, 4))
> cv.glm(Data, fit.glm.4)$delta[1]
[1] 1.114772
```

The results above are identical to the results obtained in (c) since LOOCV evaluates  $n$  folds of a single observation.

(e) Which of the models in (c) had the smallest LOOCV error? Is this what you expected? Explain your answer.

Ans:

**We may see that the LOOCV estimate for the test MSE is minimum for "fit.glm.2", this is not surprising since we saw clearly in (b) that the relation between "x" and "y" is quadratic.**

(f) Comment on the statistical significance of the coefficient estimates that results from fitting each of the models in (c) using least squares. Do these results agree with the conclusions drawn based on the cross-validation results?

Ans:

```
>summary
```

```

> summary(fit.glm.4)

Call:
glm(formula = y ~ poly(x, 4))

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.8914  -0.5244   0.0749   0.5932   2.7796

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -1.8277     0.1041  -17.549  <2e-16 ***
poly(x, 4)1    2.3164     1.0415   2.224   0.0285 *
poly(x, 4)2  -21.0586     1.0415  -20.220  <2e-16 ***
poly(x, 4)3   -0.3048     1.0415   -0.293   0.7704
poly(x, 4)4   -0.4926     1.0415   -0.473   0.6373
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 1.084654)

    Null deviance: 552.21  on 99  degrees of freedom
Residual deviance: 103.04  on 95  degrees of freedom
AIC: 298.78

Number of Fisher Scoring iterations: 2

```

The p-values show that the linear and quadratic terms are statistically significant and that the cubic and 4th degree terms are not statistically significant. This agrees strongly with our cross-validation results which were minimum for the quadratic model.





